

I will use call `IAudioFlinger::setMode` API as a sample scenario to show how the Android IPC system works. `AudioFlinger` is a service in the `media_server` program

## Service Manager Run

`service_manager` provide service manager service to other process. It must be started before any other service gets running.

```
int main(int argc, char **argv)
{
    struct binder_state *bs;
    void *svcmgr = BINDER_SERVICE_MANAGER;

    bs = binder_open(128*1024);

    if (binder_become_context_manager(bs)) {
        LOGE("cannot become context manager (%s)\n", strerror(errno));
        return -1;
    }

    svcmgr_handle = svcmgr;
    binder_loop(bs, svcmgr_handler);
    return 0;
}
```

It first open “/dev/binder” driver and then call `BINDER_SET_CONTEXT_MGR` ioctl to let binder kernel driver know it acts as a manager. Then it enters into a loop to wait for any data from other process.

```
void binder_loop(struct binder_state *bs, binder_handler func)
{
    int res;
    struct binder_write_read bwr;
    unsigned readbuf[32];

    bwr.write_size = 0;
    bwr.write_consumed = 0;
    bwr.write_buffer = 0;

    readbuf[0] = BC_ENTER_LOOPER;
    binder_write(bs, readbuf, sizeof(unsigned));

    for (;;) {
        bwr.read_size = sizeof(readbuf);
        bwr.read_consumed = 0;
        bwr.read_buffer = (unsigned) readbuf;
```

```

        res = ioctl(bs->fd, BINDER_WRITE_READ, &bwr);

        if (res < 0) {
            LOGE("binder_loop: ioctl failed (%s)\n", strerror(errno));
            break;
        }

        res = binder_parse(bs, 0, readbuf, bwr.read_consumed, func);
        if (res == 0) {
            LOGE("binder_loop: unexpected reply?!\\n");
            break;
        }
        if (res < 0) {
            LOGE("binder_loop: io error %d %s\\n", res, strerror(errno));
            break;
        }
    }
}

```

Pay attention to BINDER\_SERVICE\_MANAGER.

/\* the one magic object \*/

```
#define BINDER_SERVICE_MANAGER ((void*) 0)
```

BINDER\_SERVICE\_MANAGER is the registered handle for service\_manager. The other process must use this handle to talk with service\_manager.

## Get IServiceManager

To get an IServiceManager instance the only method is to call defaultServiceManager implemented in IServiceManager.cpp.

```

sp<IServiceManager> defaultServiceManager()
{
    if (gDefaultServiceManager != NULL) return gDefaultServiceManager;

    {
        AutoMutex _l(gDefaultServiceManagerLock);
        if (gDefaultServiceManager == NULL) {
            gDefaultServiceManager = interface_cast<IServiceManager>(
                ProcessState::self()->getContextObject(NULL));
        }
    }

    return gDefaultServiceManager;
}

```

gDefaultServiceManager is defined in libutil, so any program or library which included libutil will have this symbol, it's unique in one process. The first time gDefaultServiceManager is NULL, so it will first get a ProcessState instance through ProcessState::self(). One process has only one ProcessState instance. ProcessState will open “/dev/binder” driver for IPCThreadState use.

```
ProcessState::ProcessState()
    : mDriverFD(open_driver())
```

Now we have an instance of ProcessState, let's look at the getContextObject.

```
sp<IBinder> ProcessState::getContextObject(const sp<IBinder>& caller)
{
    if (supportsProcesses()) {
        return getStrongProxyForHandle(0);
    } else {
        return getContextObject(String16("default"), caller);
    }
}
```

The board support binder driver, so we will get into getStrongProxyForHandle. (Handle 0 is reserved for service manager, which will be explained later.)

```
sp<IBinder> ProcessState::getStrongProxyForHandle(int32_t handle)
{
    sp<IBinder> result;

    AutoMutex _l(mLock);

    handle_entry* e = lookupHandleLocked(handle);

    if (e != NULL) {
        // We need to create a new BpBinder if there isn't currently one, OR we
        // are unable to acquire a weak reference on this current one. See comment
        // in getWeakProxyForHandle() for more info about this.
        IBinder* b = e->binder;
        if (b == NULL || !e->refs->attemptIncWeak(this)) {
            b = new BpBinder(handle);
            e->binder = b;
            if (b) e->refs = b->getWeakRefs();
            result = b;
        } else {
            // This little bit of nastiness is to allow us to add a primary
            // reference to the remote proxy when this team doesn't have one
            // but another team is sending the handle to us.
            result->force_set(b);
            e->refs->decWeak(this);
        }
    }
}
```

```

    return result;
}

```

The first time b will be NULL, so the code will new an BpBinder instance. BpBinder is a base proxy class for remote binder object.

```

BpBinder::BpBinder(int32_t handle)
    : mHandle(handle)
    , mAlive(1)
    , mObitsSent(0)
    , mObituaries(NULL)
{
    LOGV("Creating BpBinder %p handle %d\n", this, mHandle);

    extendObjectLifetime(OBJECT_LIFETIME_WEAK);
    IPCThreadState::self()->incWeakHandle(handle);
}

```

IPCThreadState::incWeakHandle will add a BC\_INCREFS command in output buffer.

```

void IPCThreadState::incWeakHandle(int32_t handle)
{
    LOG_REMOTEREFS("IPCThreadState::incWeakHandle(%d)\n", handle);
    mOut.writeInt32(BC_INCREFS);
    mOut.writeInt32(handle);
}

```

Now getContextObject returns a BpBinder instance, it will be interface\_cast to IServiceManager. interface\_cast is defined in IInterface.h. It will be extended as:

```

inline sp<IServiceManager> interface_cast(const sp<IBinder>& obj)
{
    return IServiceManager::asInterface(obj);
}

```

Now let's take a look at definition of IServiceManager.

```

class IServiceManager : public IInterface
{
public:
    DECLARE_META_INTERFACE(ServiceManager);

    /**
     * Retrieve an existing service, blocking for a few seconds
     * if it doesn't yet exist.
     */
    virtual sp<IBinder>          getService( const String16& name) const = 0;

    /**
     * Retrieve an existing service, non-blocking.
     */
    virtual sp<IBinder>          checkService( const String16& name) const = 0;
}

```

```

/**
 * Register a service.
 */
virtual status_t          addService( const String16& name,
                                      const sp<IBinder>& service) = 0;

/**
 * Return list of all existing services.
 */
virtual Vector<String16>   listServices() = 0;

enum {
    GET_SERVICE_TRANSACTION = IBinder::FIRST_CALL_TRANSACTION,
    CHECK_SERVICE_TRANSACTION,
    ADD_SERVICE_TRANSACTION,
    LIST_SERVICES_TRANSACTION,
};
};

```

DECLARE\_META\_INTERFACE macro is defined as follows in IInterface.h. And it will be extended to the following code:

```

static const String16 descriptor;
static sp<IServiceManager> asInterface(const sp<IBinder>& obj);
virtual String16 getInterfaceDescriptor() const;

```

As you can see, DECLARE\_META\_INTERFACE macro declares two functions which are implemented by IMPLEMENT\_META\_INTERFACE macro in IServiceManager.cpp.

```
IMPLEMENT_META_INTERFACE(ServiceManager, "android.os.IServiceManager");
```

The code will be extended like this.

```

const String16 IServiceManager::descriptor(NAME);
String16 IServiceManager::getInterfaceDescriptor() const {
    return IServiceManager::descriptor;
}
sp<IServiceManager> IServiceManager::asInterface(const sp<IBinder>& obj)
{
    sp<IServiceManager> intr;
    if (obj != NULL) {
        intr = static_cast<IServiceManager*>(
            obj->queryLocalInterface(
                IServiceManager::descriptor).get());
        if (intr == NULL) {
            intr = new BpServiceManager(obj);
        }
    }
    return intr;
}

```

```
}
```

So `IServiceManager::asInterface` will finally new a `BpServiceManager` instance and return it to user. `BpServiceManager` works as a proxy for remote `BnServiceManager`. Any operation on `IServiceManager` now actually is to call the corresponding virtual functions in `BpServiceManager`.

Summary:

This section gives out details in how to get a proxy object for remote object.

Assume you want to implement your own service **IFunnyTest**, you must do the following:

- Put **DECLARE\_META\_INTERFACE(FunnyTest)** macro in your interface header file.
- Put **IMPLEMENT\_META\_INTERFACE(Funnytest, "your unique name")** macro in your interface source file.
- Implement your own **BpFunnyTest** class.

## Generate AudioFlinger Service

`media_server` program will start the `AudioFlinger` service. Here is the code.

```
int main(int argc, char** argv)
{
    sp<ProcessState> proc(ProcessState::self());
    sp<IServiceManager> sm = defaultServiceManager();
    LOGI("ServiceManager: %p", sm.get());
    AudioFlinger::instantiate();
    MediaPlayerService::instantiate();
    CameraService::instantiate();
    ProcessState::self()->startThreadPool();
    IPCThreadState::self()->joinThreadPool();
}
```

`AudioFlinger` will call `IServiceManager::addService`, which is a RPC call.

```
void AudioFlinger::instantiate() {
    defaultServiceManager()->addService(
        String16("media.audio_flinger"), new AudioFlinger());
}
```

`AudioFlinger` inherits from `BnAudioFlinger`, which is a template `BnInterface` class.

```
class BnAudioFlinger : public BnInterface<IAudioFlinger>
{
public:
    virtual status_t    onTransact( uint32_t code,
                                    const Parcel& data,
                                    Parcel* reply,
                                    uint32_t flags = 0);
};
```

`BnInterface` inherits from `BBinder`.

```
template<typename INTERFACE>
```

```

class BnInterface : public INTERFACE, public BBinder
{
public:
    virtual sp<IInterface>      queryLocalInterface(const String16& _descriptor);
    virtual String16            getInterfaceDescriptor() const;

protected:
    virtual IBinder*            onAsBinder();
};
template<typename INTERFACE>
IBinder* BnInterface<INTERFACE>::onAsBinder()
{
    return this;
}

```

According to BnInterface's implementation we know that the parameter passed to IServiceManager::addService is the address of new AudioFlinger instance. BBinder derives from IBinder, its transact function is to call virtual function onTransact.

```

status_t BBinder::transact(
    uint32_t code, const Parcel& data, Parcel* reply, uint32_t flags)
{
    data.setDataPosition(0);

    status_t err = NO_ERROR;
    switch (code) {
        case PING_TRANSACTION:
            reply->writeInt32(pingBinder());
            break;
        default:
            err = onTransact(code, data, reply, flags);
            break;
    }

    if (reply != NULL) {
        reply->setDataPosition(0);
    }

    return err;
}

```

The most important virtual functions is onTransact. BnAudioFlinger implemented the virtual function. For the scenario, we just need to focus on SET\_MODE branch.

```

status_t BnAudioFlinger::onTransact(
    uint32_t code, const Parcel& data, Parcel* reply, uint32_t flags)
{
    switch(code) {

```

```

    case SET_MODE: {
        CHECK_INTERFACE(IAudioFlinger, data, reply);
        int mode = data.readInt32();
        reply->writeInt32( setMode(mode) );
        return NO_ERROR;
    } break;

```

media\_server will enter into a loop through IPCThreadState::joinThreadPool, just like service\_manager, it will be waited data from other process in talkWithDriver.

```

void IPCThreadState::joinThreadPool(bool isMain)
{
    mOut.writeInt32(isMain ? BC_ENTER_LOOPER : BC_REGISTER_LOOPER);

    status_t result;
    do {
        int32_t cmd;
        result = talkWithDriver();
        if (result >= NO_ERROR) {
            size_t IN = mIn.dataAvail();
            if (IN < sizeof(int32_t)) continue;
            cmd = mIn.readInt32();
            result = executeCommand(cmd);
        }

        // Let this thread exit the thread pool if it is no longer
        // needed and it is not the main process thread.
        if(result == TIMED_OUT && !isMain) {
            break;
        }
    } while (result != -ECONNREFUSED && result != -EBADF);

    mOut.writeInt32(BC_EXIT_LOOPER);
    talkWithDriver(false);
}

```

Assume you want to implement your own service **IFunnyTest**, you must do the following:

- Implement your own **BnFunnyTest** class.
- In the process your service running, call IPCThreadState::joinThreadPool to start the loop for binder.

## RPC Call IServiceManager::addService

We called IServiceManager::addService, which means call BpServiceManager::addService.

```

virtual status_t addService(const String16& name, const sp<IBinder>& service)

```



```

{
    Parcel data, reply;
    data.writeInterfaceToken(IServiceManager::getInterfaceDescriptor());
    data.writeString16(name);
    data.writeStrongBinder(service);
    status_t err = remote()->transact(ADD_SERVICE_TRANSACTION, data, &reply);
    return err == NO_ERROR ? reply.readInt32() : err;
}

```

Parcel is simple. We can think it as a continuous buffer. Pay attention here, service parameter points to BBinder object(AudioFlinger derived from Bn)

```

status_t Parcel::writeStrongBinder(const sp<IBinder>& val)
{
    return flatten_binder(ProcessState::self(), val, this);
}

```

flatten\_binder will generate a Binder command. Because BBinder is a local binder object, so the code branch to the lines marked red.

```

status_t flatten_binder(const sp<ProcessState>& proc,
    const sp<IBinder>& binder, Parcel* out)
{
    flat_binder_object obj;

    obj.flags = 0x7f | FLAT_BINDER_FLAG_ACCEPTS_FDS;
    if (binder != NULL) {
        IBinder *local = binder->localBinder();
        if (!local) {
            BpBinder *proxy = binder->remoteBinder();
            if (proxy == NULL) {
                LOGE("null proxy");
            }
            const int32_t handle = proxy ? proxy->handle() : 0;
            obj.type = BINDER_TYPE_HANDLE;
            obj.handle = handle;
            obj.cookie = NULL;
        } else {
            obj.type = BINDER_TYPE_BINDER;
            obj.binder = local->getWeakRefs();
            obj.cookie = local;
        }
    }
    return finish_flatten_binder(binder, obj, out);
}

```

Pay attention to the red lines, local's address is put into the packet(It will be used later). After the packet for addService RPC call is made, BpServiceManager::addService will call BpBinder's transact.

```

status_t BpBinder::transact(
    uint32_t code, const Parcel& data, Parcel* reply, uint32_t flags)
{
    // Once a binder has died, it will never come back to life.
    if (mAlive) {
        status_t status = IPCThreadState::self()->transact(
            mHandle, code, data, reply, flags);
        if (status == DEAD_OBJECT) mAlive = 0;
        return status;
    }

    return DEAD_OBJECT;
}

```

BpBinder calls IPCThreadState::transact to start a transaction to the binder object corresponding to mHandle. For this scenario, mHandle is 0.

```

status_t IPCThreadState::transact(int32_t handle,
                                uint32_t code, const Parcel& data,
                                Parcel* reply, uint32_t flags)
{
    status_t err = data.errorCheck();

    flags |= TF_ACCEPT_FDS;

    if (err == NO_ERROR) {
        LOG_ONeway(">>>> SEND from pid %d uid %d %s", getpid(), getuid(),
            (flags & TF_ONE_WAY) == 0 ? "READ REPLY" : "ONE WAY");
        err = writeTransactionData(BC_TRANSACTION, flags, handle, code, data, NULL);
    }

    if (err != NO_ERROR) {
        if (reply) reply->setError(err);
        return (mLastError = err);
    }

    if ((flags & TF_ONE_WAY) == 0) {
        if (reply) {
            err = waitForResponse(reply);
        } else {
            Parcel fakeReply;
            err = waitForResponse(&fakeReply);
        }
    } else {
        err = waitForResponse(NULL, NULL);
    }
}

```

```

        return err;
    }

```

IPCThreadState::transact will first call writeTransactionData to compose a transaction structure for binder kernel driver. Pay attention to the following lines, it's very important for binder kernel driver to identify the transaction target.

```

status_t IPCThreadState::writeTransactionData(int32_t cmd, uint32_t binderFlags,
    int32_t handle, uint32_t code, const Parcel& data, status_t* statusBuffer)
{
    binder_transaction_data tr;

    tr.target.handle = handle;
    tr.code = code;
    tr.flags = binderFlags;

    const status_t err = data.errorCheck();
    if (err == NO_ERROR) {
        tr.data_size = data.ipcDataSize();
        tr.data.ptr.buffer = data.ipcData();
        tr.offsets_size = data.ipcObjectsCount()*sizeof(size_t);
        tr.data.ptr.offsets = data.ipcObjects();
    } else if (statusBuffer) {
        tr.flags |= TF_STATUS_CODE;
        *statusBuffer = err;
        tr.data_size = sizeof(status_t);
        tr.data.ptr.buffer = statusBuffer;
        tr.offsets_size = 0;
        tr.data.ptr.offsets = NULL;
    } else {
        return (mLastError = err);
    }

    mOut.writeInt32(cmd);
    mOut.write(&tr, sizeof(tr));

    return NO_ERROR;
}

```

Then waitForResponse will call talkWithDriver to invoke BINDER\_WRITE\_READ ioctl.

```

#ifdef HAVE_ANDROID_OS
    if (ioctl(mProcess->mDriverFD, BINDER_WRITE_READ, &bwr) >= 0)
        err = NO_ERROR;
    else
        err = -errno;
#else

```

So until now the transaction data has been delivered to binder kernel driver.

Summary:

Proxy object will generate the needed packet for a RPC call and then invoke BINDER\_WRITE\_READ to write the packet to binder kernel driver. The packet is a formatted packet. For RPC call, it uses BC\_TRANSACTION packet type.

Assume you want to implement your own service **IFunnyTest**, you must do the following:

- In the process your service running, call IServiceManager::addService to register the service to servier\_manger.

## Transaction in Binder Kernel Driver

When any process opens the “/dev/binder” driver, a corresponding structure binder\_proc will be assigned at binder\_open.

```
static int binder_open(struct inode *nodp, struct file *filp)
{
    struct binder_proc *proc;

    proc = kzalloc(sizeof(*proc), GFP_KERNEL);
    if (proc == NULL)
        return -ENOMEM;
    get_task_struct(current);
    proc->tsk = current;
    INIT_LIST_HEAD(&proc->todo);
    init_waitqueue_head(&proc->wait);
    proc->default_priority = task_nice(current);
    mutex_lock(&binder_lock);
    binder_stats.obj_created[BINDER_STAT_PROC]++;
    hlist_add_head(&proc->proc_node, &binder_procs);
    proc->pid = current->group_leader->pid;
    INIT_LIST_HEAD(&proc->delivered_death);
    filp->private_data = proc;
    mutex_unlock(&binder_lock);

    if (binder_proc_dir_entry_proc) {
        char strbuf[11];
        snprintf(strbuf, sizeof(strbuf), "%u", proc->pid);
        create_proc_read_entry(strbuf, S_IRUGO, binder_proc_dir_entry_proc,
binder_read_proc_proc, proc);
    }

    return 0;
}
```

So that when any ioctl reached, the driver can know the process info. The transaction data is

transferred through BINDER\_WRITE\_READ ioctl.

```
case BINDER_WRITE_READ: {
    struct binder_write_read bwr;
    if (size != sizeof(struct binder_write_read)) {
        ret = -EINVAL;
        goto err;
    }
    if (copy_from_user(&bwr, ubuf, sizeof(bwr))) {
        ret = -EFAULT;
        goto err;
    }
    if (bwr.write_size > 0) {
        ret = binder_thread_write(proc, thread, (void __user *)bwr.write_buffer,
bwr.write_size, &bwr.write_consumed);
        if (ret < 0) {
            bwr.read_consumed = 0;
            if (copy_to_user(ubuf, &bwr, sizeof(bwr)))
                ret = -EFAULT;
            goto err;
        }
    }
    if (bwr.read_size > 0) {
        ret = binder_thread_read(proc, thread, (void __user *)bwr.read_buffer,
bwr.read_size, &bwr.read_consumed, filp->f_flags & O_NONBLOCK);
        if (!list_empty(&proc->todo))
            wake_up_interruptible(&proc->wait);
        if (ret < 0) {
            if (copy_to_user(ubuf, &bwr, sizeof(bwr)))
                ret = -EFAULT;
            goto err;
        }
    }
    if (copy_to_user(ubuf, &bwr, sizeof(bwr))) {
        ret = -EFAULT;
        goto err;
    }
    break;
}
```

The driver first handles write, then read. Let's take a look at the binder\_thread\_write first. The core in binder\_thread\_write is a loop to parse command from write buffer and execute the corresponding command.

```
uint32_t cmd;
void __user *ptr = buffer + *consumed;
void __user *end = buffer + size;
```

```

while (ptr < end && thread->return_error == BR_OK) {
    if (get_user(cmd, (uint32_t __user *)ptr))
        return -EFAULT;
    ptr += sizeof(uint32_t);
    if (_IOC_NR(cmd) < ARRAY_SIZE(binder_stats.bc)) {
        binder_stats.bc[_IOC_NR(cmd)]++;
        proc->stats.bc[_IOC_NR(cmd)]++;
        thread->stats.bc[_IOC_NR(cmd)]++;
    }
    switch (cmd) {
    case ***:

    default:
        printk(KERN_ERR "binder: %d:%d unknown command %d\n", proc->pid,
thread->pid, cmd);
        return -EINVAL;
    }
    *consumed = ptr - buffer;
}

```

We just take a look at 2 commands related to the scenario. One is BC\_INCREFS.

```

case BC_INCREFS:
case BC_ACQUIRE:
case BC_RELEASE:
case BC_DECREFS: {
    uint32_t target;
    struct binder_ref *ref;
    const char *debug_string;

    if (get_user(target, (uint32_t __user *)ptr))
        return -EFAULT;
    ptr += sizeof(uint32_t);
    if (target == 0 && binder_context_mgr_node &&
        (cmd == BC_INCREFS || cmd == BC_ACQUIRE)) {
        ref = binder_get_ref_for_node(proc,
            binder_context_mgr_node);
    } else
        ref = binder_get_ref(proc, target);
    if (ref == NULL) {
        binder_user_error("binder: %d:%d refcou"
            "nt change on invalid ref %d\n",
            proc->pid, thread->pid, target);
        break;
    }
}

```

```

switch (cmd) {
case BC_INCREFS:
    debug_string = "IncRefs";
    binder_inc_ref(ref, 0, NULL);
    break;
}
break;

```

Remember that, we mentioned it before, in this scenario, the target will be 0. binder\_context\_mgr\_node is created to represent handle 0 when system\_manager calls BINDER\_SET\_CONTEXT\_MGR ioctl. So here just increase a weak reference to binder\_context\_mgr\_node node.

```
binder_context_mgr_node = binder_new_node(proc, NULL);
```

The other is BC\_TRANSACTION.

```

case BC_TRANSACTION:
case BC_REPLY: {
    struct binder_transaction_data tr;

    if (copy_from_user(&tr, ptr, sizeof(tr)))
        return -EFAULT;
    ptr += sizeof(tr);
    binder_transaction(proc, thread, &tr, cmd == BC_REPLY);
    break;
}

```

binder\_transaction will create a new binder node if the packet contains a BINDER\_TYPE\_BINDER flattened object.

```

fp = (struct flat_binder_object *) (t->buffer->data + *offp);
switch (fp->type) {
case BINDER_TYPE_BINDER:
case BINDER_TYPE_WEAK_BINDER: {
    struct binder_ref *ref;
    struct binder_node *node = binder_get_node(proc, fp->binder);
    if (node == NULL) {
        node = binder_new_node(proc, fp->binder);
        if (node == NULL) {
            return_error = BR_FAILED_REPLY;
            goto err_binder_new_node_failed;
        }
        node->cookie = fp->cookie;
    }
    ref = binder_get_ref_for_node(target_proc, node);
    if (ref == NULL) {
        return_error = BR_FAILED_REPLY;
        goto err_binder_get_ref_for_node_failed;
    }
}
}

```

```

    }
    if (fp->type == BINDER_TYPE_BINDER)
        fp->type = BINDER_TYPE_HANDLE;
    else
        fp->type = BINDER_TYPE_WEAK_HANDLE;
    fp->handle = ref->desc;
    binder_inc_ref(ref, fp->type == BINDER_TYPE_HANDLE, &thread->todo);
    if (binder_debug_mask & BINDER_DEBUG_TRANSACTION)
        printk(KERN_INFO "          node %d u%p -> ref %d desc %d\n",
                node->debug_id, node->ptr, ref->debug_id, ref->desc);
} break;

```

binder\_transaction will know that the target is handle 0, so it runs the following branch to get the target\_node, target\_proc and target\_thread.

```

    } else {
        target_node = binder_context_mgr_node;
        if (target_node == NULL) {
            return_error = BR_DEAD_REPLY;
            goto err_no_context_mgr_node;
        }
    }
    e->to_node = target_node->debug_id;
    target_proc = target_node->proc;
    if (!(tr->flags & TF_ONE_WAY) && thread->transaction_stack) {
        struct binder_transaction *tmp;
        tmp = thread->transaction_stack;
        while (tmp) {
            if (tmp->from && tmp->from->proc == target_proc)
                target_thread = tmp->from;
            tmp = tmp->from_parent;
        }
    }
}

```

Finally binder\_transaction will put this request into list and wake up the waiting thread in binder\_thread\_read.

```

t->work.type = BINDER_WORK_TRANSACTION;
list_add_tail(&t->work.entry, target_list);
tcomplete->type = BINDER_WORK_TRANSACTION_COMPLETE;
list_add_tail(&tcomplete->entry, &thread->todo);
if (target_wait)
    wake_up_interruptible(target_wait);

```

Now let's take a look at binder\_thread\_read. When service\_manager runs, it will wait at here until a request is delivered to it.

```

ret = wait_event_interruptible_exclusive(proc->wait, binder_has_proc_work(proc,
thread));

```



Because the previous write from media\_server process has wakened it up, so it gets executed. The following code is to copy data from write buffer of media\_server to read buffer of system\_manager.

```
tr.data_size = t->buffer->data_size;
tr.offsets_size = t->buffer->offsets_size;
tr.data.ptr.buffer = (void *)((void *)t->buffer->data + proc->user_buffer_offset);
tr.data.ptr.offsets = tr.data.ptr.buffer + ALIGN(t->buffer->data_size, sizeof(void *));

if (put_user(cmd, (uint32_t __user *)ptr))
    return -EFAULT;
ptr += sizeof(uint32_t);
if (copy_to_user(ptr, &tr, sizeof(tr)))
    return -EFAULT;
ptr += sizeof(tr);
```

Summary:

This section has shown the data flow from RPC call client side to RPC server side.

## Service Manager Handle Add Service

Until now, the service\_manager has gotten a packet of BR\_TRANSACTION from media\_server, so it will call binder\_parser to handle the packet.

```
case BR_TRANSACTION: {
    if (func) {
        unsigned rdata[256/4];
        struct binder_io msg;
        struct binder_io reply;
        int res;

        bio_init(&reply, rdata, sizeof(rdata), 4);
        bio_init_from_txn(&msg, txn);
        res = func(bs, txn, &msg, &reply);
        binder_send_reply(bs, &reply, txn->data, res);
    }
    ptr += sizeof(*txn) / sizeof(uint32_t);
    break;
}
```

binder\_parser will call svcmgr\_handler to parse BR\_TRANSACTION packet, which is the reverse process of BpServerManager. Here structure binder\_txn actually is the same with structure binder\_transaction\_data. In our scenario, the transaction code is SVC\_MGR\_ADD\_SERVICE.

```
int svcmgr_handler(struct binder_state *bs,
                  struct binder_txn *txn,
                  struct binder_io *msg,
```

```

                                struct binder_io *reply)
{
    struct svcinfo *si;
    uint16_t *s;
    unsigned len;
    void *ptr;

    if (txn->target != svcmgr_handle)
        return -1;

    s = bio_get_string16(msg, &len);

    if ((len != (sizeof(svcmgr_id) / 2)) ||
        memcmp(svcmgr_id, s, sizeof(svcmgr_id))) {
        fprintf(stderr, "invalid id %s\n", str8(s));
        return -1;
    }

    switch(txn->code) {
    case SVC_MGR_ADD_SERVICE:
        s = bio_get_string16(msg, &len);
        ptr = bio_get_ref(msg);
        if (do_add_service(bs, s, len, ptr, txn->sender_euid))
            return -1;
        break;

```

So that service\_manager knows that a service named s will run and it gets the object information through bio\_get\_ref.

```

void *bio_get_ref(struct binder_io *bio)
{
    struct binder_object *obj;

    obj = _bio_get_obj(bio);
    if (!obj)
        return 0;

    if (obj->type == BINDER_TYPE_HANDLE)
        return obj->pointer;

    return 0;
}

```

bio\_get\_ref is the reverse work of flatten\_binder. do\_add\_service will finally call BC\_ACQUIRE to get a strong reference to the object represented by ptr.

Summary:

This section shows how the service is added to service manager.

Assume you want to implement your own service **IFunnyTest**, you must do the following:

- Add your service name into allowed service list in service\_manager.

## Get IAudioFlinger

The only way to get a service interface is through `IServiceManager::getService`. For example, here is the method for `AudioSystem` to get an `IAudioFlinger`.

```
// establish binder interface to AudioFlinger service
const sp<IAudioFlinger>& AudioSystem::get_audio_flinger()
{
    Mutex::Autolock _l(gLock);
    if (gAudioFlinger.get() == 0) {
        sp<IServiceManager> sm = defaultServiceManager();
        sp<IBinder> binder;
        do {
            binder = sm->getService(String16("media.audio_flinger"));
            if (binder != 0)
                break;
            LOGW("AudioFlinger not published, waiting...");
            usleep(500000); // 0.5 s
        } while(true);
        gAudioFlinger = interface_cast<IAudioFlinger>(binder);
    }
    LOGE_IF(gAudioFlinger==0, "no AudioFlinger!?");
    return gAudioFlinger;
}
```

`IServiceManager::getService` is called to `BpServiceManager::getService`.

```
virtual sp<IBinder> getService(const String16& name) const
{
    unsigned n;
    for (n = 0; n < 5; n++){
        sp<IBinder> svc = checkService(name);
        if (svc != NULL) return svc;
        LOGI("Waiting for service %s...\n", String8(name).string());
        sleep(1);
    }
    return NULL;
}

virtual sp<IBinder> checkService( const String16& name) const
{
    Parcel data, reply;
    data.writeInterfaceToken(IServiceManager::getInterfaceDescriptor());
```

```

        data.writeString16(name);
        remote()->transact(CHECK_SERVICE_TRANSACTION, data, &reply);
        return reply.readStrongBinder();
    }

```

Just like analyzed before, the call will finally through Binder kernel driver to be handled in service\_manager process.

```

switch(txn->code) {
case SVC_MGR_GET_SERVICE:
case SVC_MGR_CHECK_SERVICE:
    s = bio_get_string16(msg, &len);
    ptr = do_find_service(bs, s, len);
    if (!ptr)
        break;
    bio_put_ref(reply, ptr);
    return 0;
}

```

Then service\_manager will reply with the previous handle set by media\_server (Which is the AudioFlinger's instance's address.). Then BpServiceManager::checkService will be returned back from remote()->transact call. Then just like analyzed for IServiceManager, it will new another BpBinder instance corresponding to the handle returned back from service\_manager. interface\_cast<IAudioFlinger>(binder) will finally return an BpAudioFlinger instance.

Summary:

Just like get IServiceManager, but this time it needs to get a handle from service\_manager. While for IServiceManager it always use handle 0.

## RPC Call IAudioFlinger::SetMode

If we call IAudioFlinger::SetMode in AAA process, actually call BpAudioFlinger::setMode.

```

virtual status_t setMode(int mode)
{
    Parcel data, reply;
    data.writeInterfaceToken(IAudioFlinger::getInterfaceDescriptor());
    data.writeInt32(mode);
    remote()->transact(SET_MODE, data, &reply);
    return reply.readInt32();
}

```

Like analysis of IServiceManager::addService, this function will generate a packet and write it to binder kernel driver, then wait to read reply. The only difference is this time the target handle points to some address in media\_server process.

## Handle IAudioFlinger::SetMode

The Binder kernel driver will finally wake up the read thread of media\_server process which runs at IPCThreadState::joinThreadPool. Now let's re-look the code.

```
void IPCThreadState::joinThreadPool(bool isMain)
{
    mOut.writeInt32(isMain ? BC_ENTER_LOOPER : BC_REGISTER_LOOPER);

    status_t result;
    do {
        int32_t cmd;
        result = talkWithDriver();
        if (result >= NO_ERROR) {
            size_t IN = mIn.dataAvail();
            if (IN < sizeof(int32_t)) continue;
            cmd = mIn.readInt32();
            result = executeCommand(cmd);
        }

        // Let this thread exit the thread pool if it is no longer
        // needed and it is not the main process thread.
        if (result == TIMED_OUT && !isMain) {
            break;
        }
    } while (result != -ECONNREFUSED && result != -EBADF);

    mOut.writeInt32(BC_EXIT_LOOPER);
    talkWithDriver(false);
}
```

At this time, talkWithDriver will return with the packet generated by BpServiceManager::setMode, then executeCommand will process the commands. In this scenario, the command is BR\_TRANSACTION.

```
case BR_TRANSACTION:
{
    binder_transaction_data tr;

    Parcel reply;
    if (tr.target.ptr) {
        sp<BBinder> b((BBinder*)tr.cookie);
        const status_t error = b->transact(tr.code, buffer, &reply, 0);
        if (error < NO_ERROR) reply.setError(error);
    } else {
        const status_t error = the_context_object->transact(tr.code, buffer, &reply, 0);
    }
}
```

```

        if (error < NO_ERROR) reply.setError(error);
    }

    if ((tr.flags & TF_ONE_WAY) == 0) {
        LOG_ONeway("Sending reply to %d!", mCallingPid);
        sendReply(reply, 0);
    } else {
        LOG_ONeway("NOT sending reply to %d!", mCallingPid);
    }
}

break;

```

The most important two lines are marked with red. Here it gets an address from binder kernel driver and cast it to a BBinder pointer(This address is put into binder kernel driver when called IServiceManager::addService.). Remember that, AudioFlinger is derived from BBinder. This pointer is actually the same pointer of our AudioFlinger instance. So the following transact call will finally call our BnAudioFlinger's onTransact virtual function.

```

status_t BnAudioFlinger::onTransact(
    uint32_t code, const Parcel& data, Parcel* reply, uint32_t flags)
{
    case SET_MODE: {
        CHECK_INTERFACE(IAudioFlinger, data, reply);
        int mode = data.readInt32();
        reply->writeInt32( setMode(mode) );
        return NO_ERROR;
    } break;
}

```

Then the reply will be written through sendReply.

```


status_t IPCThreadState::sendReply(const Parcel& reply, uint32_t flags)
{
    status_t err;
    status_t statusBuffer;
    err = writeTransactionData(BC_REPLY, flags, -1, 0, reply, &statusBuffer);
    if (err < NO_ERROR) return err;

    return waitForResponse(NULL, NULL);
}

```

It will finally write to binder kernel driver. Then kernel driver will wake up the read thread of AAA.

## Summary



The overall architecture of Android IPC system is shown in the picture. There are four major blocks:

- **Binder Driver**  
It's the core of IPC system. It passes data between service provider and service user.
- **service provider**  
It provides some kind of service. It will parser the received RPC call data from binder driver and do the real action.
- **service\_manager**  
It's a special service provider. It provides service manager service for other service provider.
- **service user**  
It remote calls service provider. It will generate an RPC call data and send it to binder driver.

For the example scenario, here lists the major control flow.

1. **service\_manager** runs first, it will register a special node 0 in binder driver.
2. **media\_server** gets an **IServiceManager** proxy object for the special node 0.
3. **media\_server** RPC call **IServiceManager::addService** to add the **IAudioFlinger** service. This call is routed to node 0. It will send data to binder driver.
4. binder driver notify that data is for node 0, and the data contains command to binder a object. So it will generate another node (assume A) for the **IAudioFlinger** service and route the data to **service\_manager**.
5. **service\_manager** reads data from binder driver, then process **IServiceManager::addService** RPC call.
6. another process P gets an **IServiceManager** proxy object for the special node 0.
7. P RPC call **IServiceManager::getService** to get **IAudioFlinger** service. This call is routed to node 0. It will send data to binder driver.
8. binder driver notify that data is for node 0, So it will route the data to **service\_manager**.
9. **service\_manager** reads data from binder driver, then process **IServiceManager::getService** RPC call and return back the node A to represent the **IAudioFlinger** service.
10. P RPC call **IAudioFlinger::setMode**. Now this call will be routed to node A.
11. binder driver notify that data is for node A, so it will route the data to **media\_server**.
12. **media\_server** reads data from binder driver, handles **IAudioFlinger::setMode** RPC call and send reply data to binder driver.
13. binder driver route the reply to P.
14. P reads data from binder driver and finally gets the reply.